

GENIFY

LECTURE 20

14-September-2025

TOPICS COVERED:

RIS


- RETRIEVAL AUGMENTED GENERATION (RAG)

Presenter

Shahzaib Kashif

Co-Presenter

Shayan Hassan Baig



RETRIEVAL AUGMENTED GENERATION (RAG)



Q: What is Retrieval-Augmented Generation (RAG) ?

- Retrieval-augmented generation (RAG) is an innovative approach in the field of NLP.
- It combines the strengths of retrieval-based and generation-based models to enhance the quality of generated text.

What is RAG?

Retrieving relevant data and generating accurate, context-aware responses to improve AI outputs.

R

Retrieve

Find useful information

A

Augment

Add it to the AI's knowledge

G

Generate

Create a better response


Q: Why is Retrieval-Augmented Generation important?

- In traditional LLMs, the model generates responses based solely on the data it was trained on, which may not include the most current information or specific details required for certain tasks.
- RAG addresses this limitation by incorporating a retrieval mechanism that allows the model to access external databases or documents in real-time.

This hybrid model aims to leverage the vast amounts of information available in large-scale databases or knowledge bases making it particularly effective for tasks that require accurate and contextually relevant information.

Q: How does Retrieval-Augmented Generation work?

The system first searches external sources for relevant information based on the user's query Instead of relying only on existing training data.



1. Creating External Data

- External data refers to new information beyond the LLM's original training dataset.
- It can come from various sources, such as APIs, databases, or document repositories, and may exist in different formats like text files or structured records.
- To make this data understandable to AI, it is first divided in chunks in case of massive datasets and converted into numerical representations (embeddings) using specialized models and then stored in a vector database
- This creates a knowledge library that the AI system can reference during retrieval.

2. Retrieving Relevant Information

- When a user submits a query, the system converts it into a vector representation and matches it against stored vectors in the database.
- This enables precise retrieval of the most relevant information.
- For example, if the LLM is asked, "What are the key topics in the DSA course?"
- It would retrieve both the course syllabus and relevant study materials.
- This ensures the response is highly relevant and tailored to the user's learning needs.

3. Augmenting the LLM Prompt

- Once the relevant data is retrieved;
- it is incorporated into the user's input (prompt) using prompt engineering techniques.
- This enhances the model's contextual understanding, allowing it to generate more detailed, factually accurate, and insightful responses.

4. Keeping External Data Updated

- To ensure the system continues to provide reliable and up-to-date responses, external data must be refreshed periodically.
- This can be done through automated real-time updates or scheduled batch processing.
- Keeping vector embeddings updated allows the RAG system to always retrieve the most current and relevant information for generating responses.

What problems does RAG solve?

The retrieval-augmented generation (RAG) approach helps solve several challenges in natural language processing (NLP) and AI applications:

- **Factual Inaccuracies and Hallucinations:** Traditional generative models can produce plausible but incorrect information. RAG reduces this risk by retrieving verified, external data to ground responses in factual knowledge.
- **Outdated Information:** Static models rely on training data that may become obsolete. RAG dynamically retrieves up-to-date information, ensuring relevance and accuracy in real-time.
- **Contextual Relevance:** Generative models often struggle with maintaining context in complex or multi-turn conversations. RAG retrieves relevant documents to enrich the context, improving coherence and relevance.
- **Domain-Specific Knowledge:** Generic models may lack expertise in specialized fields. RAG integrates domain-specific external knowledge for tailored and precise responses.
- **Cost and Efficiency:** Fine-tuning large models for specific tasks is expensive. RAG eliminates the need for retraining by dynamically retrieving relevant data, reducing costs and computational load.
- **Scalability Across Domains:** RAG is adaptable to diverse industries, from healthcare to finance, without extensive retraining, making it highly scalable

Challenges and Future Directions

Despite its advantages, RAG faces several challenges:

- **Complexity:** Combining retrieval and generation adds complexity to the model, requiring careful tuning and optimization to ensure both components work seamlessly together.
- **Latency:** The retrieval step can introduce latency, making it challenging to deploy RAG models in real-time applications.
- **Quality of Retrieval:** The overall performance of RAG heavily depends on the quality of the retrieved documents. Poor retrieval can lead to suboptimal generation, undermining the model's effectiveness.
- **Bias and Fairness:** Like other AI models, RAG can inherit biases present in the training data or retrieved documents, necessitating ongoing efforts to ensure fairness and mitigate biases.

RAG Applications with Examples

1. Advanced Question-Answering System

- **Scenario:** Imagine a customer support chatbot for an online store. A customer asks, "What is the return policy for a damaged item?"
- **RAG in Action:** The chatbot retrieves the store's return policy document from its knowledge base. RAG then uses this information to generate a clear and concise answer like, "If your item is damaged upon arrival, you can return it free of charge within 30 days of purchase. Please visit our returns page for detailed instructions."

RAG Applications with Examples

2. Content Creation and Summarization

- **Scenario:** You're building a travel website and want to create a summary of the Great Barrier Reef.
- **RAG in Action:** RAG can access and process vast amounts of information about the Great Barrier Reef from various sources. It can then provide a concise summary highlighting key points like its location, size, biodiversity, and conservation efforts.

RAG Applications with Examples

3. Conversational Agents and Chatbots

- **Scenario:** A virtual assistant for a financial institution. A user asks, "What are some factors to consider when choosing a retirement plan?"
- **RAG in Action:** The virtual assistant retrieves relevant information about retirement plans and investment strategies. RAG then uses this knowledge to provide the user with personalized guidance based on their age, income, and risk tolerance.

RAG Applications with Examples

3. Conversational Agents and Chatbots

- **Scenario:** A virtual assistant for a financial institution. A user asks, "What are some factors to consider when choosing a retirement plan?"
- **RAG in Action:** The virtual assistant retrieves relevant information about retirement plans and investment strategies. RAG then uses this knowledge to provide the user with personalized guidance based on their age, income, and risk tolerance.

RAG Applications with Examples

4. Information Retrieval

- **Scenario:** You are searching the web for information about the history of artificial intelligence (AI).
- **RAG in Action:** A RAG-powered search engine can not only return relevant webpages but also generate informative snippets that summarize the content of each page. This allows you to quickly grasp the key points of each result without having to visit every single webpage.

RAG Applications with Examples

5. Educational Tools and Resources

- **Scenario:** An online learning platform for science courses. A student is studying about the human body and has a question about the function of the heart.
- **RAG in Action:** The platform uses RAG to access relevant information about the heart's anatomy and function from the course materials. It then presents the student with an explanation, diagrams, and perhaps even links to video resources, all tailored to their specific learning needs.